

# Times, Clocks, and the Ordering of Events in a Distributed System

Jacopo Notarstefano

`jacopo.notarstefano [at] gmail.com`

13 February 2013

# Main objectives and outline

In this paper Lamport:

- ① Discusses the partial ordering defined by the “happened before” relation,
- ② Gives a distributed algorithm for extending it to a consistent total ordering of all the events,
- ③ Uses this algorithm to solve a mutual exclusion problem.

# Distributed system

## Definition (Distributed system)

A *distributed system* consists of a collection of distinct processes which are spatially separated, and which communicate with one another by exchanging messages. A system is distributed if the message transmission delay is not negligible compared to the time between events in a single process.

# Examples of distributed systems

- A worldwide network of interconnected computers
- A cluster of workstation in a data center
- Processes on a single computer

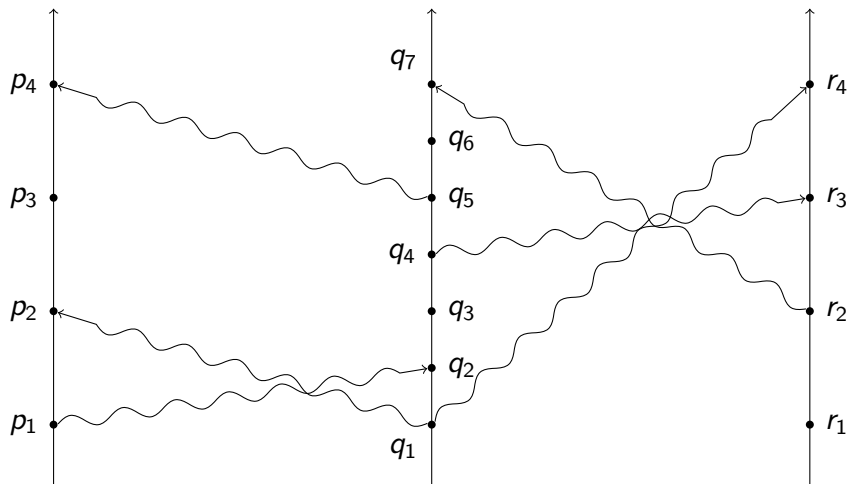
# The “ $\rightarrow$ ” relation

## Definition (The “ $\rightarrow$ ” relation)

The “ $\rightarrow$ ” relation on the set of events of a system is the smallest relation satisfying the following three conditions:

- 1 If  $a$  and  $b$  are events in the same process, and  $a$  comes before  $b$ , then  $a \rightarrow b$ .
- 2 If  $a$  is the sending of a message by one process, and  $b$  is the receipt of the same message by another process, then  $a \rightarrow b$ .
- 3 If  $a \rightarrow b$ , and  $b \rightarrow c$ , then  $a \rightarrow c$ .

# The “space-time diagram”



## Definition (Clock)

For each process  $P_i$  we define a *clock*  $C_i$  to be a function that assigns a number  $C_i\langle a \rangle$  to each event  $a$  in the process.

## Definition (System of clocks)

A *system of clocks* is a function  $C$  that assigns to the event  $b$  in process  $P_j$  the time  $C\langle b \rangle = C_j\langle b \rangle$ .

# The clock condition

## Definition (The clock condition)

We say that a system of clocks satisfies the *clock condition* if, for any events  $a$  and  $b$ , we have: if  $a \rightarrow b$  then  $C\langle a \rangle < C\langle b \rangle$ .

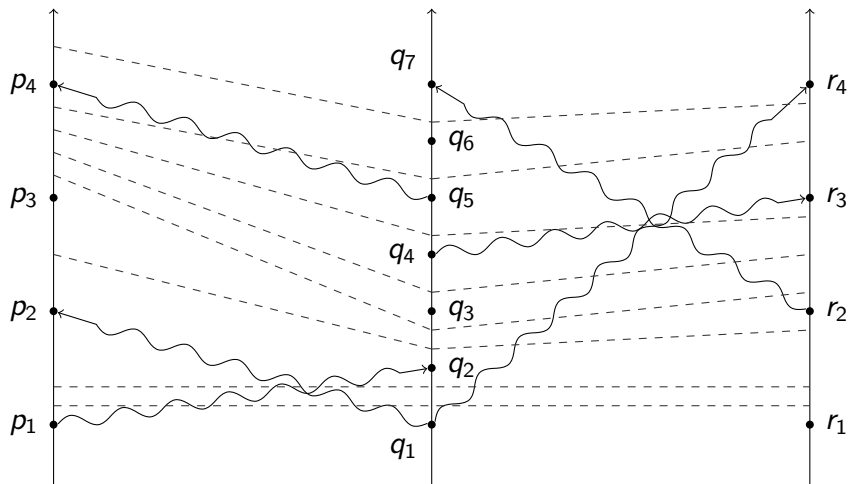
## Lemma

*The clock condition is satisfied if the following conditions hold:*

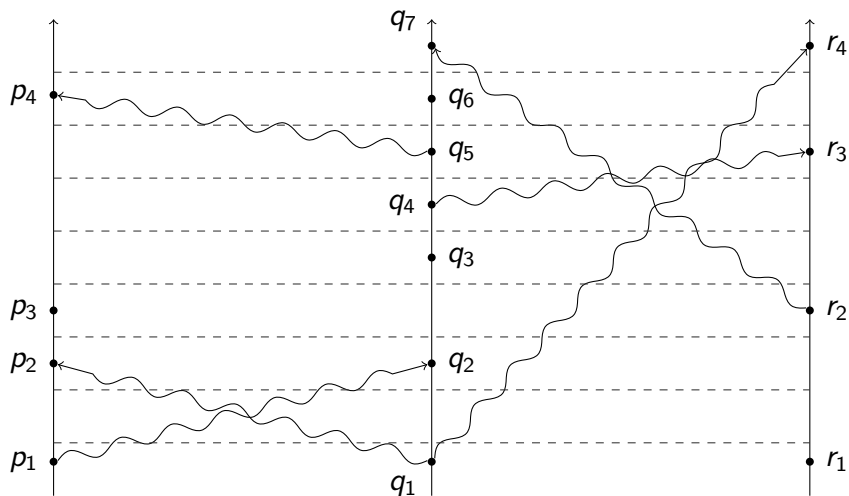
- 1 If  $a$  and  $b$  are events in process  $P_i$  and  $a$  comes before  $b$ , then  $C_i\langle a \rangle < C_i\langle b \rangle$ .
- 2 If  $a$  is the sending of a message by process  $P_i$  and  $b$  is the receipt of that message by process  $P_j$ , then  $C_i\langle a \rangle < C_j\langle b \rangle$ .



# The “space-time diagram”, revisited



# The “space-time diagram”, rearranged



# Implementation of the clock condition

## Lemma

*To guarantee that the system of clocks satisfies the clock condition we need to obey the following implementation rules:*

- 1 *Each process  $P_i$  increments  $C_i$  between any two successive events.*
- 2 *If event  $a$  is the sending of a message  $m$  by process  $P_i$ , then the message  $m$  contains a timestamp  $T_m = C_i\langle a \rangle$ .*
- 3 *Upon receiving a message  $m$ , process  $P_j$  sets  $C_j$  greater than or equal to its present value and greater than  $T_m$ .*

# The “ $\Rightarrow$ ” relation

## Definition (The “ $\Rightarrow$ ” relation)

Let  $\prec$  be a total ordering on the processes. If  $a$  is an event in process  $P_i$  and  $b$  is an event in process  $P_j$ , then  $a \Rightarrow b$  if and only if either

- 1  $C_i\langle a \rangle < C_j\langle b \rangle$  or
- 2  $C_i\langle a \rangle = C_j\langle b \rangle$  and  $P_i \prec P_j$ .

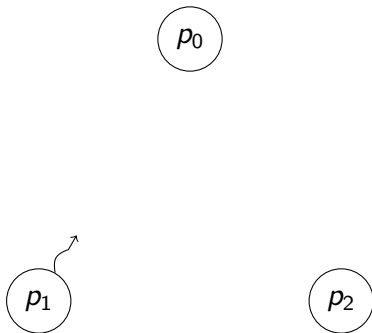
# A mutual exclusion problem

A fixed collection of processes share a single resource, which can be used by one process at a time. We want to find an algorithm that satisfies the following conditions:

- 1 A process which has been granted the resource must release it before it can be granted to another process.
- 2 Different requests for the resource must be granted in the order in which they are made.
- 3 If every process which is granted the resource eventually releases it, then every request is eventually granted.

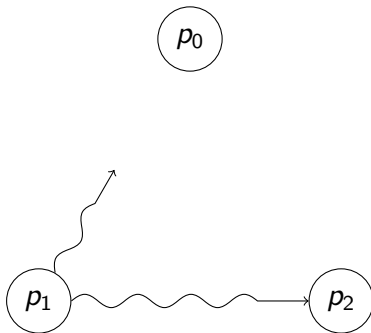
# A wrong solution, 1/3

We can't use a central scheduling process which grants requests in the order they are received. Suppose that  $p_1$  sends a message to  $p_0$ , the scheduler, to request the resource.



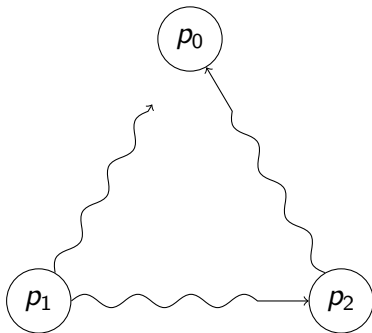
## A wrong solution, 2/3

Then  $p_1$  sends a message to  $p_2$ . Suppose that this message is received while the other is still traveling due to network congestion.



## A wrong solution, 3/3

Finally,  $p_2$  sends a message to  $p_0$  requesting the resource.  $p_0$  then grants the resource to  $p_2$ , which requested the resource after  $p_1$ .





# Some assumptions

To simplify the problem, we make some assumptions.

- ① For any two processes  $P_i$  and  $P_j$ , the messages sent from  $P_i$  to  $P_j$  are received in the same order as they are sent.
- ② Every message is eventually received.
- ③ A process can send messages directly to every other process.

# The request queue

Let  $P_0$  be the process to which the shared resource is initially allocated. Let  $T_0$  be less than the initial value of any logical clock in the system. Each process maintains a private *request queue*, which initially contains one message: " $T_0$ :  $P_0$  requests resource".

# Resource request and acknowledgment

- 1 Process  $P_i$  sends the message " $T_m$ :  $P_i$  requests resource" to every other process where  $T_m$  is the process clock's value at the time of the request.  $P_i$  also puts the request message on its request queue.
- 2 When process  $P_j$  receives  $P_i$ 's request message, it puts it in its queue and sends an acknowledgment message to  $P_i$ .

# Resource release and acknowledgement

- ③ Process  $P_i$  removes request message " $T_m: P_i$  requests resource" from its queue and sends the release message " $P_i$  releases resource" to every other process.
- ④ Process  $P_j$  removes the " $T_m: P_i$  requests resource" from its queue.

# Resource allocation

- ⑤ Process  $P_i$  is allocated the resource when:
  - There is a “ $T_m$ :  $P_i$  requests” resource in  $P_i$ 's request queue which is before any other request according to  $\Rightarrow$ .
  - $P_i$  has received messages from every other process timestamped later than  $T_m$ .

# Proof of correctness (sketch)

## Theorem

*The algorithm described by rules 1 to 5 is a correct solution to the mutual exclusion problem.*

## Proof (sketch).

The ordering is guaranteed by the fact that relation " $\Rightarrow$ " extends relation " $\rightarrow$ ".

