

Times, Clocks, and the Ordering of Events in a Distributed System

Jacopo Notarstefano

`jacopo.notarstefano [at] gmail.com`

13 February 2013

Main objectives and outline

In this paper Lamport:

- ① Discusses the partial ordering defined by the “happened before” relation,
- ② Gives a distributed algorithm for extending it to a consistent total ordering of all the events,
- ③ Uses this algorithm to solve a mutual exclusion problem.

Distributed system

Definition (Distributed system)

A *distributed system* consists of a collection of distinct processes which are spatially separated, and which communicate with one another by exchanging messages. A system is distributed if the message transmission delay is not negligible compared to the time between events in a single process.

Examples of distributed systems

- A worldwide network of interconnected computers
- A cluster of workstation in a data center
- Processes on a single computer

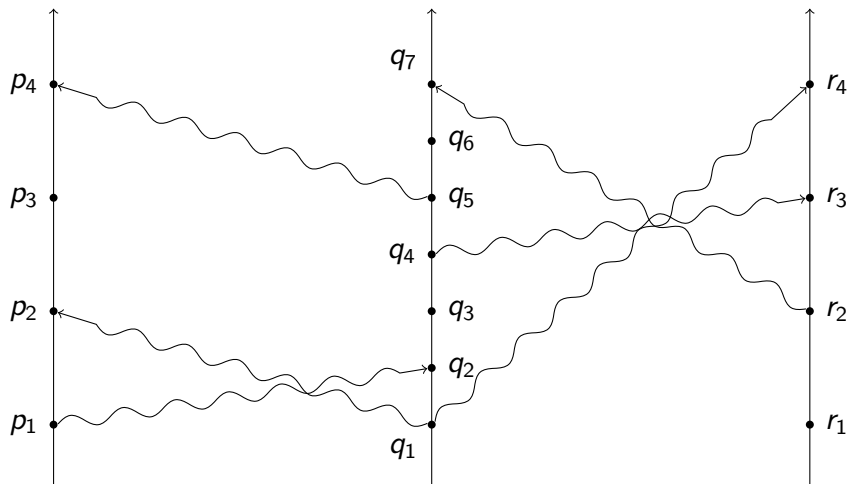
The “ \rightarrow ” relation

Definition (The “ \rightarrow ” relation)

The “ \rightarrow ” relation on the set of events of a system is the smallest relation satisfying the following three conditions:

- 1 If a and b are events in the same process, and a comes before b , then $a \rightarrow b$.
- 2 If a is the sending of a message by one process, and b is the receipt of the same message by another process, then $a \rightarrow b$.
- 3 If $a \rightarrow b$, and $b \rightarrow c$, then $a \rightarrow c$.

The “space-time diagram”



Clocks

Definition (Clock)

For each process P_i we define a *clock* C_i to be a function that assigns a number $C_i\langle a \rangle$ to each event a in the process.

Definition (System of clocks)

A *system of clocks* is a function C that assigns to the event b in process P_j the time $C\langle b \rangle = C_j\langle b \rangle$.

The clock condition

Definition (The clock condition)

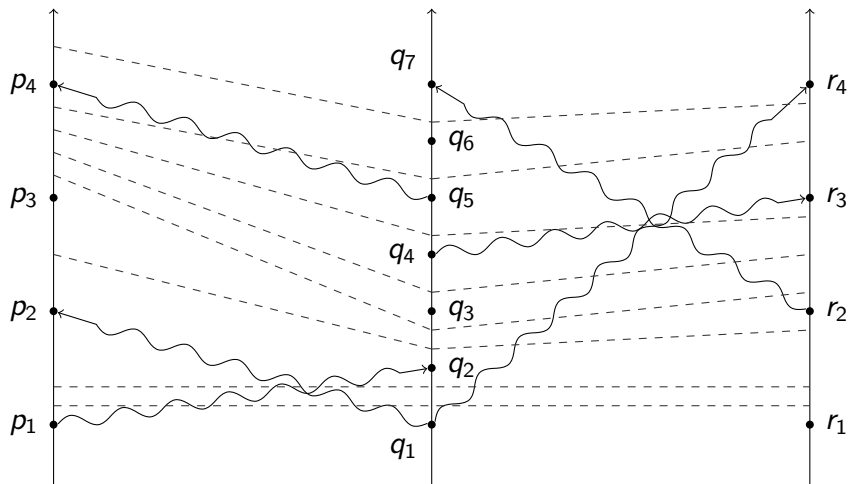
We say that a system of clocks satisfies the *clock condition* if, for any events a and b , we have: if $a \rightarrow b$ then $C\langle a \rangle < C\langle b \rangle$.

Lemma

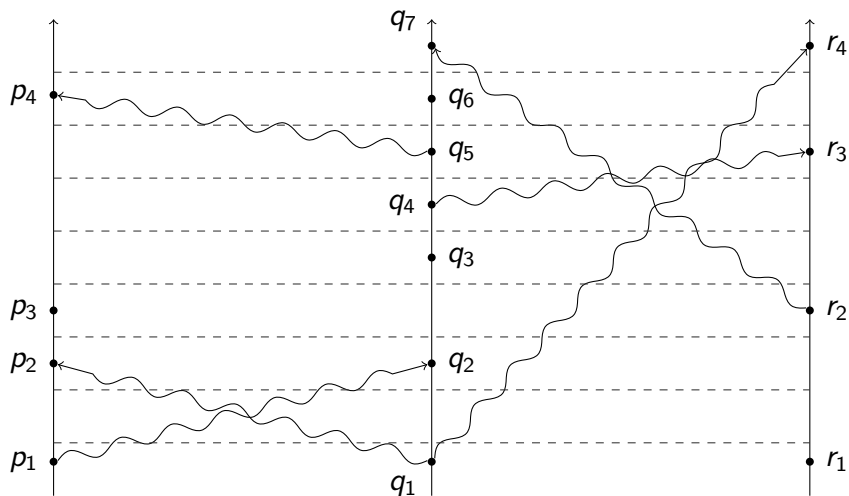
The clock condition is satisfied if the following conditions hold:

- 1 If a and b are events in process P_i and a comes before b , then $C_i\langle a \rangle < C_i\langle b \rangle$.
- 2 If a is the sending of a message by process P_i and b is the receipt of that message by process P_j , then $C_i\langle a \rangle < C_j\langle b \rangle$.

The “space-time diagram”, revisited



The “space-time diagram”, rearranged



Implementation of the clock condition

Lemma

To guarantee that the system of clocks satisfies the clock condition we need to obey the following implementation rules:

- 1 *Each process P_i increments C_i between any two successive events.*
- 2 *If event a is the sending of a message m by process P_i , then the message m contains a timestamp $T_m = C_i\langle a \rangle$.*
- 3 *Upon receiving a message m , process P_j sets C_j greater than or equal to its present value and greater than T_m .*

The “ \Rightarrow ” relation

Definition (The “ \Rightarrow ” relation)

Let \prec be a total ordering on the processes. If a is an event in process P_i and b is an event in process P_j , then $a \Rightarrow b$ if and only if either

- 1 $C_i\langle a \rangle < C_j\langle b \rangle$ or
- 2 $C_i\langle a \rangle = C_j\langle b \rangle$ and $P_i \prec P_j$.


A mutual exclusion problem

A fixed collection of processes share a single resource, which can be used by one process at a time. We want to find an algorithm that satisfies the following conditions:

- 1 A process which has been granted the resource must release it before it can be granted to another process.
- 2 Different requests for the resource must be granted in the order in which they are made.
- 3 If every process which is granted the resource eventually releases it, then every request is eventually granted.


A wrong solution, 1/3

We can't use a central scheduling process which grants requests in the order they are received.



A wrong solution, 2/3

We can't use a central scheduling process which grants requests in the order they are received.



A wrong solution, 3/3

We can't use a central scheduling process which grants requests in the order they are received.

